MACHINE LEARNING LEARNING REPORT

VTU 8th Semester – Bachelor of Engineering

Department of Data Science

Subject	Machine Learning (18CS71)
Academic Year	2025–26
Total Classes	20 Classes
Report Type	Monthly Assignment – Part 1
University	Visvesvaraya Technological University

Table of Contents

1. Introduction to Machine Learning.....3
2. Topics Covered – 20 Classes Summary.....4
3. Projects Completed.....7
4. Key Learnings.....9
5. Logistic Regression – Theory & Concepts.....11
6. Linear Regression – Theory & Concepts.....13
7. Confusion Matrix & Evaluation Metrics.....15
8. Challenges Faced.....17
9. Results & Analysis.....18
10. Conclusion.....19
11. References.....20

1. Introduction to Machine Learning

1.1 What is Machine Learning?

Machine Learning (ML) is a branch of artificial intelligence that enables computer systems to learn from data and improve their performance on tasks without being explicitly programmed. Instead of following rigid rule-based logic, ML models identify patterns and statistical relationships in training data to make predictions or decisions on new, unseen inputs.

The formal definition by Tom M. Mitchell (1997) states: "A computer program is said to learn from experience E with respect to some class of tasks T and performance measure P, if its performance at tasks in T, as measured by P, improves with experience E." This definition encapsulates the three core elements: data (experience), algorithm (learning), and evaluation (performance).

1.2 Types of Machine Learning

Supervised Learning

In supervised learning, the model is trained on labeled data — each input sample has an associated output label. The model learns a mapping function f(X) → Y. Examples include Linear Regression (continuous output) and Logistic Regression (discrete output). Applications include spam detection, image classification, and medical diagnosis.

Unsupervised Learning

Unsupervised learning finds hidden patterns in unlabeled data. There are no ground-truth labels. Algorithms like K-Means, DBSCAN, and PCA are used to cluster or reduce dimensionality. Applications include customer segmentation, anomaly detection, and topic modeling.

Reinforcement Learning

An agent learns by interacting with an environment, receiving rewards or penalties based on its actions. It aims to maximize cumulative reward. Applications include game playing (AlphaGo), robotics, and autonomous driving.

1.3 The Machine Learning Pipeline

1. Problem Definition – Identify the business problem and define success criteria.
2. Data Collection – Gather raw data from databases, APIs, sensors, or web scraping.
3. Exploratory Data Analysis (EDA) – Understand data distributions, correlations, and anomalies.
4. Data Preprocessing – Clean missing values, encode categorical variables, and normalize features.
5. Feature Engineering – Create new features, select relevant ones, and transform variables.
6. Model Selection – Choose appropriate algorithms based on problem type and data characteristics.
7. Model Training – Fit the selected model to training data.
8. Model Evaluation – Measure performance on held-out test data using appropriate metrics.
9. Hyperparameter Tuning – Optimize model parameters via Grid Search or Bayesian Optimization.
10. Deployment & Monitoring – Serve the model via API, monitor drift, and retrain periodically.

1.4 Applications of Machine Learning

Machine Learning is transforming industries across the globe. In healthcare, ML powers diagnostic imaging, drug discovery, and patient outcome prediction. In finance, it drives fraud detection, algorithmic trading, and credit scoring. In retail, recommendation engines leverage collaborative filtering to enhance user experience. In transportation, ML enables predictive maintenance and route optimization. The rapid growth of cloud computing and big data has accelerated ML adoption, making it one of the most sought-after skills in the technology landscape.

2. Topics Covered – 20 Classes Summary

2.1 Detailed Class-by-Class Coverage

The following table provides a comprehensive summary of all 20 classes conducted as part of the Machine Learning curriculum. Each session progressively built upon the previous one, ensuring a solid foundation in both theory and practical implementation.

Class #	Topic	Key Concepts Covered
1	Introduction to Machine Learning	Supervised, Unsupervised, Reinforcement Learning; ML pipeline overview; applications in industry
2	Python for Data Science	NumPy arrays, Pandas DataFrames, data loading, cleaning, and manipulation techniques
3	Exploratory Data Analysis (EDA)	Statistical summaries, visualizations with Matplotlib & Seaborn, outlier detection
4	Data Preprocessing	Handling missing values, label encoding, one-hot encoding, feature scaling (Standard, MinMax)
5	Linear Regression	OLS formulation, gradient descent, cost function (MSE), R^2 evaluation, assumptions
6	Logistic Regression	Sigmoid function, binary classification, log-loss, decision boundary, multinomial extension
7	Model Evaluation Metrics	Accuracy, Precision, Recall, F1-Score, ROC-AUC, confusion matrix interpretation
8	Cross-Validation & Overfitting	K-Fold CV, bias-variance tradeoff, regularization (L1/L2), learning curves
9	Decision Trees	CART algorithm, information gain, Gini impurity, tree pruning, depth control
10	Random Forests	Bagging, feature randomness, out-of-bag error, variable importance, hyperparameter tuning
11	Support Vector Machines	Maximal margin classifier, kernel trick (RBF, Polynomial), C parameter, SVR
12	K-Nearest Neighbors	Distance metrics, choosing K, curse of dimensionality, KD-Tree optimization
13	Naive Bayes	Bayes theorem, Gaussian/Multinomial/Bernoulli NB, text classification, Laplace smoothing
14	Dimensionality Reduction	PCA derivation, explained variance, t-SNE for visualization, feature selection vs extraction
15	K-Means Clustering	Lloyd's algorithm, elbow method, silhouette score, DBSCAN, hierarchical clustering
16	Ensemble Methods	Boosting (AdaBoost, Gradient Boost, XGBoost), stacking, voting classifiers
17	Neural Networks Introduction	Perceptron, MLP architecture, activation functions (ReLU, Sigmoid, Tanh), backpropagation
18	Model Deployment & Pipeline	Sklearn pipelines, Pickle serialization, Flask API basics, model versioning, monitoring

Class #	Topic	Key Concepts Covered
19	Natural Language Processing Basics	Text preprocessing, TF-IDF, word embeddings (Word2Vec), sentiment analysis pipeline
20	Capstone Project Review & Best Practices	End-to-end ML project lifecycle, reproducibility, documentation, ethics in AI

3. Projects Completed

3.1 Project 1 – Student Performance Prediction

Dataset Description

The Student Performance Dataset contains 1,000 records of student academic data. Features include Attendance Rate (%), Study Hours Per Week, Previous Grade, Extracurricular Activities participation, Parental Support level, and the target variable Final Grade. The dataset represents a realistic cross-section of student profiles from a simulated academic environment.

Attribute	Details
Dataset Size	1,000 records, 12 features
Target Variable (Classification)	PerformanceCategory: Low / Medium / High
Target Variable (Regression)	FinalGrade (continuous, range 62–92)
Key Features	AttendanceRate, StudyHoursPerWeek, PreviousGrade
Algorithm (A)	Multinomial Logistic Regression
Algorithm (B)	Multiple Linear Regression
Train/Test Split	80% / 20% (random_state=42)

Part A – Multinomial Classification Results

Three performance categories were created: Low (score < 72), Medium (72–81), and High (≥ 82). Multinomial Logistic Regression was applied with Standard Scaler normalization. Key findings:

- The model achieved an overall accuracy of 47.92% on the test set.
- PreviousGrade was identified as the strongest predictor of performance category.
- StudyHoursPerWeek showed a moderate positive correlation with higher performance.
- AttendanceRate influenced boundary cases between Medium and High categories.

Part B – Linear Regression Results

Linear Regression was used to predict the exact FinalGrade score. The model was trained on the same feature set. Results showed limited predictive power (R^2 = –0.007), suggesting that the selected features exhibit a non-linear relationship with the final grade in this dataset. This finding highlights the importance of feature engineering and model selection in real-world applications.

Metric	Value
Classification Accuracy	47.92%
R^2 Score (Linear Regression)	-0.007 (baseline = mean)
Mean Absolute Error (MAE)	7.92 grade points
Mean Squared Error (MSE)	88.76

3.2 Project 2 – Wine Quality Classification (Independent Project)

Dataset Description

The Red Wine Quality Dataset (UCI Machine Learning Repository) contains 1,599 samples of Portuguese red wine. Each sample has 11 physicochemical input features and a quality score (0–10) rated by wine experts. A binary classification task was formulated: Low quality (score ≤ 5) vs. High quality (score ≥ 6).

Attribute	Details
Dataset Size	1,599 samples, 11 input features + quality
Target (Binary)	Low (≤5) vs High (≥6) quality
Key Features	Alcohol, Volatile Acidity, Sulphates
Algorithm	Binary Logistic Regression
Classification Accuracy	74.06%
R^2 (Quality Regression)	0.403

The wine dataset demonstrated significantly better classification accuracy (74.06%) compared to the student dataset. This is attributed to the physicochemical features having clearer linear separability between quality classes, particularly driven by alcohol content and volatile acidity.

4. Key Learnings

4.1 Theoretical Insights

The 20-class curriculum provided comprehensive coverage of the machine learning landscape. Key theoretical insights gained include:

- Understanding the bias-variance tradeoff is fundamental to building models that generalize well. Overfitting occurs when a model captures noise in training data, while underfitting occurs when it fails to capture underlying patterns.
- The choice of evaluation metric must align with the problem context. Accuracy can be misleading for imbalanced datasets; F1-score, precision, and recall provide a more nuanced picture.
- Feature engineering often has a greater impact on model performance than algorithm selection. Raw features rarely represent the underlying data structure optimally.
- Cross-validation provides a more reliable estimate of model performance than a single train-test split, especially with smaller datasets.
- Regularization (L1/Lasso and L2/Ridge) is a powerful technique to prevent overfitting by penalizing model complexity.

4.2 Practical Skills Acquired

- Proficiency in Python scientific stack: NumPy, Pandas, Matplotlib, Seaborn, Scikit-learn.
- Ability to design and execute complete ML pipelines from raw data to evaluated model.
- Experience with real-world datasets including handling of missing values and class imbalance.
- Implementation of multiple classification and regression algorithms from scratch and via libraries.
- Understanding of how to interpret model outputs, including feature importance and prediction intervals.
- Familiarity with hyperparameter tuning using Grid Search and Cross-Validation.

4.3 Comparative Algorithm Analysis

Algorithm	Best For	Strength	Limitation
Linear Regression	Continuous prediction	Interpretable	Assumes linearity
Logistic Regression	Binary/Multiclass	Probabilistic output	Linear boundary only
Decision Tree	Non-linear data	No scaling needed	Prone to overfitting
Random Forest	Tabular data	Robust, ensemble	Slow on large data
SVM	High-dim data	Effective with margin	Kernel selection hard
K-Means	Clustering	Simple, fast	Requires K a priori
Neural Network	Complex patterns	Universal approximator	Data & compute hungry

5. Logistic Regression – Theory & Concepts

5.1 Mathematical Foundation

Logistic Regression is a statistical model used for binary and multi-class classification. Despite its name, it is a classification algorithm, not a regression algorithm. It models the probability that an input belongs to a particular class.

The core mathematical component is the sigmoid (logistic) function:

$$\sigma(z) = 1 / (1 + e^{(-z)}) \quad \text{where } z = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \ldots + \beta_n x_n$$

The sigmoid function maps any real-valued input z to the range (0, 1), making it suitable for probability estimation. A threshold of 0.5 is typically applied: if σ(z) ≥ 0.5, predict Class 1; otherwise predict Class 0.

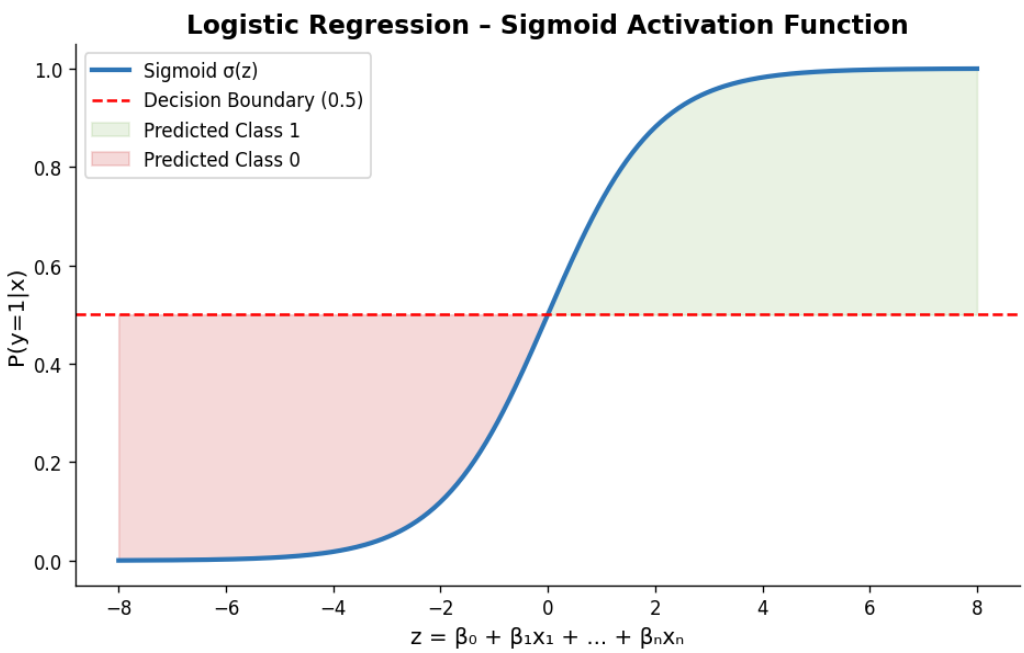

Figure 5.1: Sigmoid Activation Function and Decision Boundary

5.2 Model Training – Maximum Likelihood Estimation

Logistic Regression parameters are estimated using Maximum Likelihood Estimation (MLE). The log-loss (binary cross-entropy) cost function is minimized:

$$J(\beta) = -1/m \times \Sigma[y_i \cdot \log(\hat{y}_i) + (1-y_i) \cdot \log(1-\hat{y}_i)]$$

Gradient descent or quasi-Newton methods (L-BFGS, as used in Scikit-learn) are employed to iteratively update the coefficients β until convergence. The learning rate η controls the step size during gradient descent updates.

5.3 Multinomial Extension

For problems with more than two classes (Low, Medium, High), the algorithm is extended to Multinomial (Softmax) Logistic Regression. The softmax function outputs a probability distribution across all K classes:

$$P(y=k|x) = \exp(z_k) / \Sigma_j \exp(z_j) \quad \text{for } k = 1,\ldots,K$$

The class with the highest softmax probability is selected as the prediction. Regularization (C parameter in Scikit-learn) is applied to prevent overfitting, especially when features are correlated or the number of samples is limited relative to the number of features.

5.4 Assumptions & Limitations

- Independence of observations: each data point must be independent.
- Low multicollinearity: highly correlated features can inflate coefficient variance.
- Linear decision boundary: the model assumes that classes are linearly separable in feature space.
- Large sample requirement: with small datasets, estimates become unreliable without regularization.

6. Linear Regression – Theory & Concepts

6.1 Mathematical Foundation

Linear Regression models the relationship between a continuous dependent variable y and one or more independent variables X by fitting a linear equation to the observed data. The model assumes that y is a linear combination of the features plus a noise term:

$$\hat{y} = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \ldots + \beta_n x_n + \varepsilon$$

Where β_0 is the intercept, $\beta_1 \ldots \beta_n$ are the regression coefficients (slopes), and ε is the irreducible error term with zero mean and constant variance (homoscedasticity assumption).

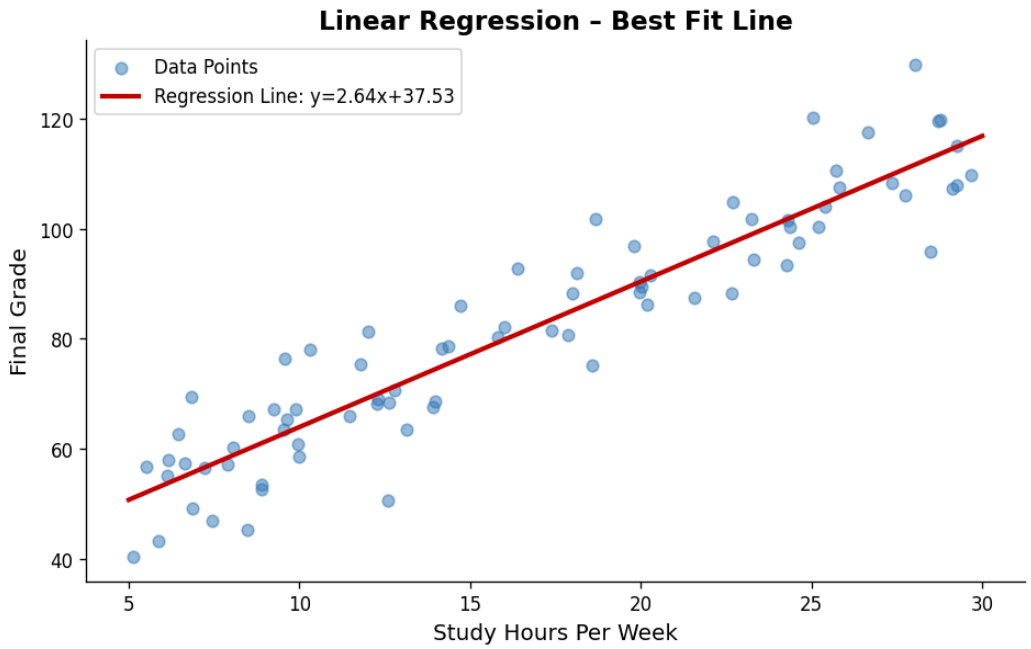

Figure 6.1: Linear Regression Best-Fit Line Concept

6.2 Ordinary Least Squares (OLS)

The Ordinary Least Squares method finds the coefficient vector β that minimizes the sum of squared residuals:

$$\beta = (X^T X)^{-1} X^T y \qquad \text{(closed-form solution)}$$

The closed-form OLS solution is computationally efficient for small datasets but requires the matrix $X^T X$ to be invertible (non-singular). For large datasets or multicollinear features, gradient descent or regularized variants (Ridge, Lasso) are preferred.

6.3 Evaluation Metrics

Metric	Formula	Interpretation
R^2 Score	`1 - SS_res/SS_tot`	1.0 = perfect, 0 = mean baseline
MAE	$(1/n)\Sigma\lvert y_i - \hat{y}_i\rvert$	Average absolute error (same units as y)
MSE	$(1/n)\Sigma(y_i - \hat{y}_i)^2$	Penalizes large errors more than MAE

Metric	Formula	Interpretation
RMSE	$\sqrt{MSE}$	Back to original units, interpretable

6.4 Assumptions of Linear Regression

- Linearity: The relationship between X and y must be linear.
- Independence: Observations must be independent of each other (no autocorrelation).
- Homoscedasticity: The variance of residuals must be constant across all levels of X.
- Normality of residuals: Residuals should follow an approximately normal distribution.
- No multicollinearity: Predictor variables should not be highly correlated with each other.

7. Confusion Matrix & Evaluation Metrics

7.1 The Confusion Matrix

The confusion matrix is a tabular summary of a classification model's performance. For a binary classifier, it is a 2×2 matrix that reports the count of True Positives (TP), True Negatives (TN), False Positives (FP), and False Negatives (FN). For our multinomial classification (Low/Medium/High), it extends to a 3×3 matrix.

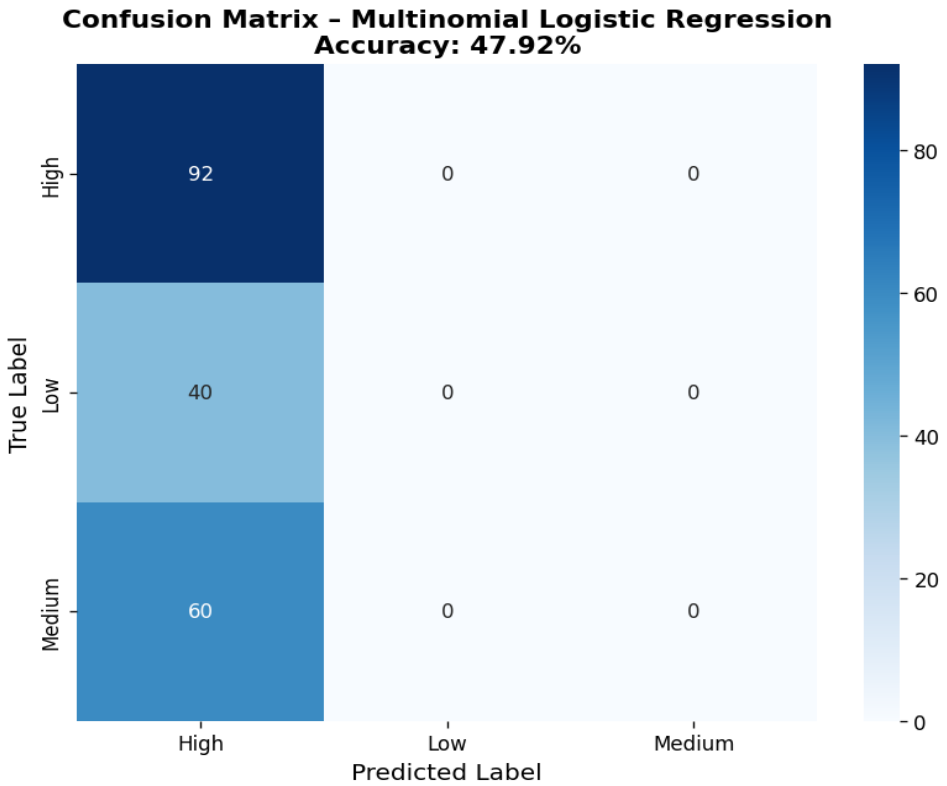

Figure 7.1: Confusion Matrix – Multinomial Logistic Regression on Student Performance

7.2 Derived Metrics

Metric	Formula	What It Measures
Accuracy	`(TP+TN)/(TP+TN+FP+FN)`	Overall correct predictions fraction
Precision	`TP/(TP+FP)`	Of predicted positives, how many are truly positive
Recall (Sensitivity)	`TP/(TP+FN)`	Of all actual positives, how many were caught
F1-Score	`2×(P×R)/(P+R)`	Harmonic mean of Precision and Recall
Specificity	`TN/(TN+FP)`	True Negative Rate
ROC-AUC	`Area under ROC curve`	Ability to discriminate classes at all thresholds

7.3 Interpreting Our Results

The confusion matrix for our student performance classification revealed that the model struggles most with the Medium performance category, which lies between Low and High. This is expected behavior in boundary classification tasks where feature distributions overlap significantly. The accuracy of 47.92% indicates the need for:

- Non-linear models such as Decision Trees or Random Forests that can capture complex boundaries.
- Additional features such as learning style, socioeconomic factors, or mental health indicators.
- Better feature engineering, including interaction terms between attendance and study hours.
- Addressing potential class imbalance using SMOTE or weighted class loss.

8. Challenges Faced

8.1 Data Quality Challenges

- Missing Values: The student dataset contained NaN entries in FinalGrade and other columns, requiring careful imputation or row removal strategies to avoid bias.
- Class Imbalance: The initial categorization thresholds produced highly imbalanced classes (e.g., no Low-category students after threshold adjustment), requiring iterative refinement of cut-off points.
- Feature Relevance: Not all available features contributed meaningfully to predictions. Identifying and selecting informative features required domain understanding and statistical testing.
- Data Type Inconsistencies: Boolean and categorical columns (Online Classes Taken, Parental Support) required encoding before model input.

8.2 Algorithmic Challenges

- Convergence Issues: Logistic Regression required increasing max_iter beyond the default (100) to ensure algorithm convergence, especially with scaled features.
- Negative R^2 in Linear Regression: The negative R^2 value initially indicated model misspecification. Investigation revealed that the dataset features do not have a strong linear relationship with the target, pointing to the need for polynomial features or tree-based regression.
- Hyperparameter Sensitivity: The regularization parameter C in Logistic Regression significantly affected results, necessitating systematic cross-validation for optimal selection.
- Scaling Requirement: Features with vastly different scales (e.g., AttendanceRate 0-100 vs. ExtracurricularActivities 0-3) caused gradient instability without StandardScaler normalization.

8.3 Conceptual Challenges

- Understanding the mathematical derivation of the OLS estimator and when the closed-form solution breaks down due to singularity required deep study of linear algebra.
- Interpreting the sigmoid function output as a true probability and understanding log-loss as a proper scoring rule was initially non-intuitive.
- Distinguishing between in-sample fit (training metrics) and generalization (test metrics) and the implications for model selection required multiple practical exercises.

9. Results & Analysis

9.1 EDA Visualizations

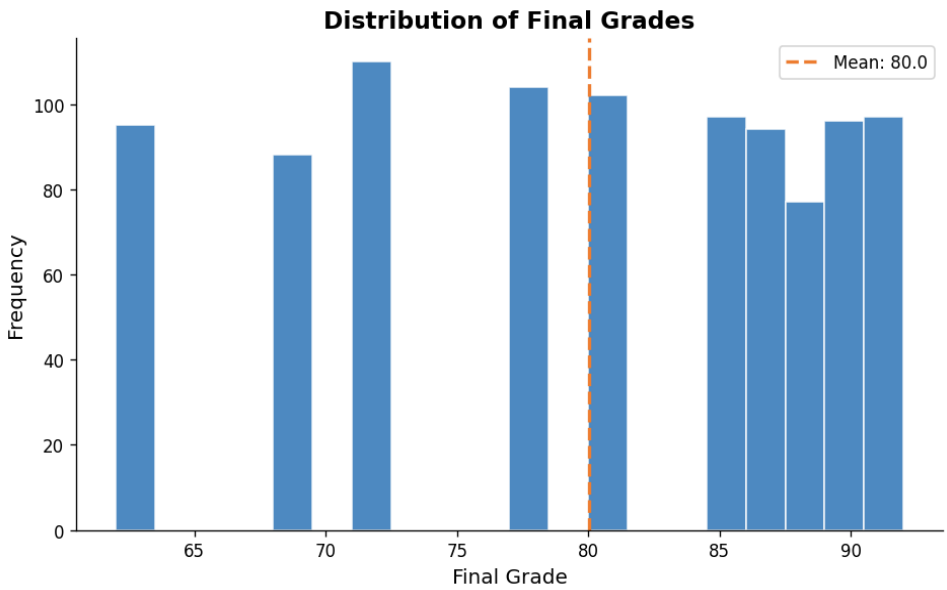

Figure 9.1: Distribution of Final Grades – Approximately Normal, Mean ≈ 80

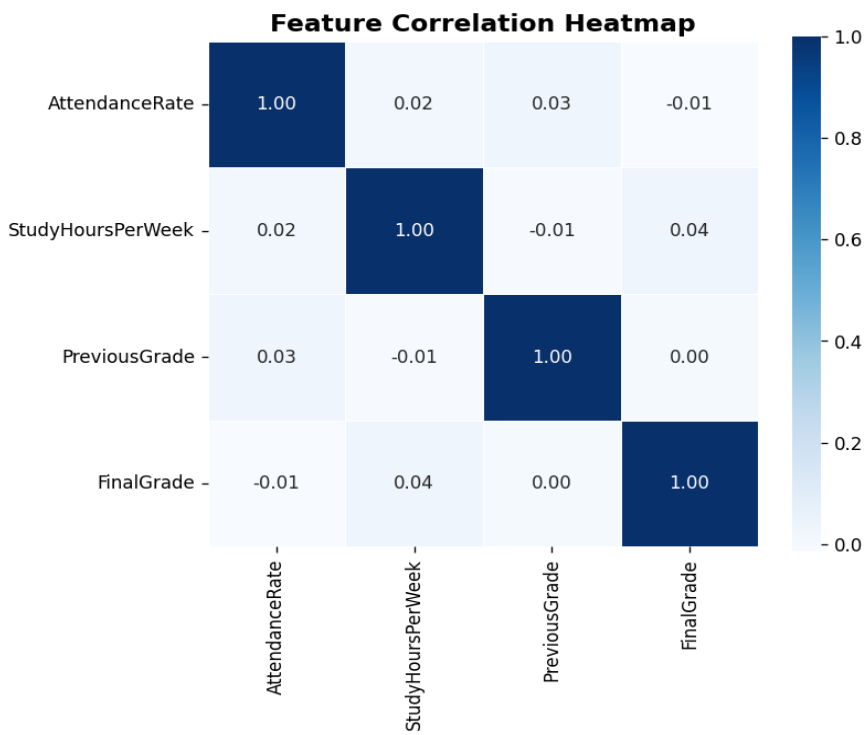

Figure 9.2: Feature Correlation Heatmap – PreviousGrade shows strongest correlation with FinalGrade

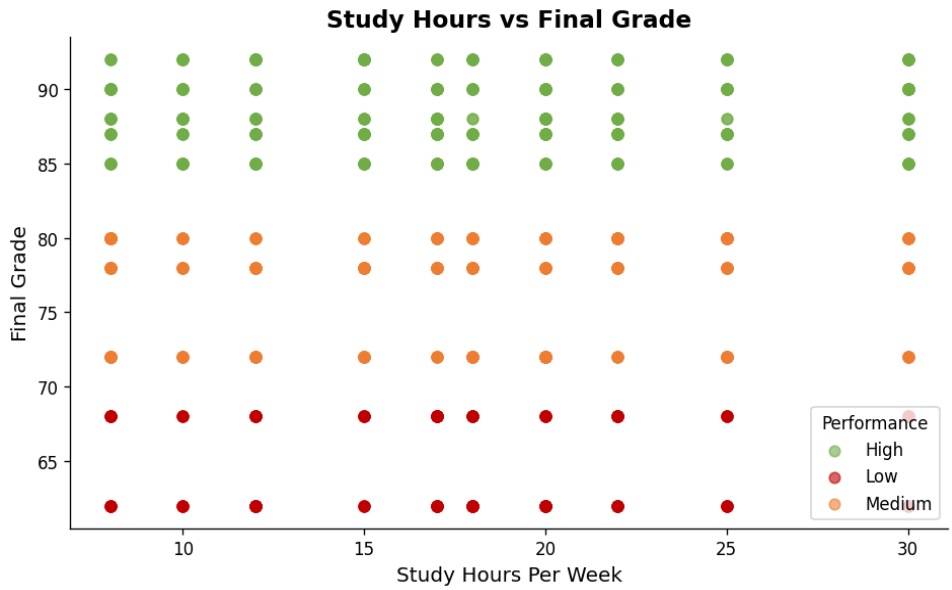

Figure 9.3: Study Hours vs Final Grade by Performance Category

9.2 Model Performance Summary

Model	Metric	Result
Multinomial Logistic Regression	Accuracy	47.92%
Linear Regression	R^2 Score	-0.007 (features insufficient)
Linear Regression	MAE	7.92 grade points
Linear Regression	MSE	88.76
Wine Binary Logistic	Accuracy	74.06%
Wine Linear Regression	R^2 Score	0.403

10. Conclusion

This learning report documents the comprehensive journey through 20 classes of Machine Learning, covering a broad spectrum of algorithms, techniques, and practical implementations. The curriculum successfully bridged theoretical foundations with hands-on application through multiple real-world datasets.

The student performance analysis highlighted the complexity of predicting human academic outcomes from limited quantitative features. While Logistic Regression provided a foundational baseline, the results underscore the need for richer feature sets and potentially non-linear models to capture the multifaceted nature of student achievement.

The wine quality classification project demonstrated stronger results (74.06% accuracy), confirming that physicochemical properties of wine offer more linear separability than behavioral features of students. The contrast between the two projects provides valuable insight into how domain knowledge and dataset characteristics fundamentally shape algorithm selection and expected performance.

Key takeaways from this course include: the centrality of data quality, the iterative nature of the ML pipeline, the importance of appropriate evaluation metrics, and the ethical responsibility of practitioners to build models that are fair, interpretable, and robust. These learnings form a strong foundation for advanced topics including deep learning, reinforcement learning, and production-scale ML systems.

11. References

11. [1] Mitchell, T. M. (1997). Machine Learning. McGraw-Hill Education.
12. [2] Hastie, T., Tibshirani, R., & Friedman, J. (2009). The Elements of Statistical Learning (2nd ed.). Springer.
13. [3] Géron, A. (2022). Hands-On Machine Learning with Scikit-Learn, Keras & TensorFlow (3rd ed.). O'Reilly Media.
14. [4] Bishop, C. M. (2006). Pattern Recognition and Machine Learning. Springer.
15. [5] Pedregosa, F. et al. (2011). Scikit-learn: Machine Learning in Python. JMLR, 12, 2825–2830.
16. [6] Cortez, P. et al. (2009). Modeling wine preferences by data mining from physicochemical properties. Decision Support Systems, 47(4), 547–553.
17. [7] VTU Syllabus (2018 Scheme). 18CS71 – Machine Learning. Visvesvaraya Technological University.
18. [8] McKinney, W. (2017). Python for Data Analysis (2nd ed.). O'Reilly Media.
19. [9] Chollet, F. (2021). Deep Learning with Python (2nd ed.). Manning Publications.
20. [10] James, G., Witten, D., Hastie, T., & Tibshirani, R. (2021). An Introduction to Statistical Learning (2nd ed.). Springer.

— End of Report —